

Improving the Practical Space and Time Efficiency of the Shortest-Paths Approach to Sum-of-Pairs Multiple Sequence Alignment

SANDEEP K. GUPTA,¹ JOHN D. KECECIOGLU,² and ALEJANDRO A. SCHÄFFER¹

ABSTRACT

The MSA program, written and distributed in 1989, is one of the few existing programs that attempts to find optimal alignments of multiple protein or DNA sequences. The MSA program implements a branch-and-bound technique together with a variant of Dijkstra's shortest paths algorithm to prune the basic dynamic programming graph. We have made substantial improvements in the time and space usage of MSA. The improvements make feasible a variety of problem instances that were not feasible previously. On some runs we achieve an order of magnitude reduction in space usage and a significant multiplicative factor speedup in running time. To explain how these improvements work, we give a much more detailed description of MSA than has been previously available. In practice, MSA rarely produces a provably optimal alignment and we explain why.

INTRODUCTION

ALIGNMENT OF MULTIPLE DNA and protein sequences is a fundamental problem in molecular biology. A variety of combinatorial definitions of the alignment problem are used in practice. Finding an optimal multiple sequence alignment seems to require time and space exponential in the number of sequences for all definitions, and is provably hard for a few (Maier, 1978; Bodlaender *et al.*, 1994; Kececioğlu, 1993; Wang and Jiang, 1994). Nevertheless it is often useful to align a small number of sequences, that is greater than two, optimally.

One of the existing programs that attempts to construct an optimal alignment, for some definition of multiple sequence alignment, is **MSA** (Carrillo and Lipman, 1988; Lipman *et al.*, 1989). **MSA** (version 1.0) was completed and distributed in 1989. Most of the implementation of the resource-intensive part was performed by J.D. Kececioğlu, and the shortest paths computation description in the sections below is derived from notes written by J.D. Kececioğlu in 1989. Lipman *et al.* (1989) give a short overview of the program, but no detailed description of the implementation has ever been published.

Some other programs that attempt to compute a global multiple sequence alignment include **AMULT** (Barton and Sternberg, 1987a,b), **DFALIGN** (Feng and Doolittle, 1987), **MULTAL** (Taylor, 1987, 1988),

¹Department of Computer Science, Rice University, Houston, TX 77005-1892.

²Department of Computer Science, University of Georgia, Athens, GA 30602-7404.

Present address of S.K. Gupta: Laboratory for Computer Science, Massachusetts Institute of Technology, 545 Technology Square, Cambridge, MA 02139; skgupta@pdos.lcs.mit.edu.

TULLA (Subbiah and Harrison, 1989), **CLUSTAL V** (Higgins *et al.*, 1992), and **MWT** (Kececiloglu, 1993). Of these methods, **MWT** is the only one other than **MSA** that attempts to compute an optimal alignment based on some well-defined score. In practice, **MSA** rarely produces a provably optimal alignment; the reasons for lack of optimality are explained at the end of the second section. Methods other than **MSA** and **MWT** tend to use much less space because they do not search for an optimal alignment. Two surveys of these programs and other issues in multiple sequence alignment have been given (Chan *et al.*, 1992; McClure *et al.*, 1994). In the second survey, **MSA** was not tested extensively, precisely because it exhausted the space on the computer that McClure *et al.* (1994) used.

The stimulus for this paper is a successful attempt to reduce the space usage of **MSA** by significant multiplicative factors. On some runs, we achieve an order of magnitude reduction in space usage and a significant multiplicative improvement in running time. The amount of improvement varies widely with the data set. In practice, the space usage of the original **MSA** is a more severe constraint than its time usage. The space reduction we have achieved makes many runs that were not feasible with the original **MSA** feasible for the first time, and speeds up many runs that before were marginally feasible but would thrash extensively while allocating virtual memory. *Thrashing* occurs when the amount of space that a process requires approaches or exceeds the available physical memory and the operating system must repeatedly *swap* memory items between the physical memory and the virtual memory on disk. Many of these thrashing runs can now fit in main memory, which can lead to a huge double-digit multiplicative reduction in wall-clock time.

Our approach to improving the **MSA** program uses algorithmic techniques that are well-known in sectors of the computer science research community, but not well-known in the biology research community. The resource-intensive part of **MSA** is an implementation of a complex variant of Dijkstra's single-source shortest-paths algorithm (Dijkstra, 1959). The version of multiple sequence alignment that **MSA** solves can be formulated as a single-source, single-destination shortest-paths problem on a mesh-shaped dynamic programming graph. The dimension of the mesh is the number of sequences, and the number of vertices is the product of the sequence lengths.

To save space (even in the original version) the dynamic programming graph is generated "on-the-fly," so that only subgraphs are stored in memory at any given time. The essential idea in our space usage improvements is to prove algorithmic invariants regarding when edges and vertices of the graph first need to be created, and when they can be safely destroyed. To further improve the running time, we carefully recoded a few subroutines that consumed the bulk of the running time, as detected by runtime profiling of the code.

Though **MSA** will never be able to compete with heuristic alignment algorithms that can align dozens of sequences of triple-digit lengths, the new version is significantly more powerful than the original. The source code is currently available by anonymous ftp from softlib.cs.rice.edu. Contact A.A. Schäffer by e-mail for details.

The paper is organized as follows. The second section defines the version of multiple sequence alignment that **MSA** solves and the basic mathematical method it uses. The third section describes in more detail how **MSA** implements this method in the absence of gap penalties. The fourth section explains how gap penalties are incorporated. The fifth section describes our space and time improvements. The sixth section compares the two versions of **MSA** with measurements on real data sets of protein sequences. We conclude with a short discussion.

DEFINITIONS AND THE CARRILLO-LIPMAN METHOD

In this section we define the version of multiple sequence alignment solved by **MSA** and review the method of Carrillo and Lipman that motivated the development of **MSA**. Let S_1, \dots, S_K , be the input sequences and assume that K is at least 2. Let Σ be the input alphabet; we assume that Σ does not contain the character '-', so that a dash can be used to denote a gap in the alignment. Most of the inputs to **MSA** consist of protein sequences, but the only place in which this assumption shows up is that the program contains default character substitution costs for proteins.

A multiple sequence alignment A is a rectangular character array on the alphabet $\Sigma' := \Sigma \cup \{-\}$ that satisfies the following conditions:

```
--VLSPADKTNVKAAWGKVGAGHAGEYGAEALE-
-VHLTPEEKSAVTALWGKVNVND--EVGGEALGR
--GLSDGEWQLVLNVWGKVEADIPGHGQEVLI-
MKFFAVLALCIVGAIASPLTADEASLVQSS---
```

FIG. 1. An alignment of prefixes of four globin sequences.

1. There are exactly K rows.
2. Ignoring dashes, row I is precisely the string S_I .

Each multiple sequence alignment induces a pairwise alignment $A_{i,j}$ on the pair of sequences S_i, S_j . This alignment is obtained by simply copying rows i and j of A , with the proviso that columns with a ‘-’ in both rows are ignored.

Figure 1 shows an example of multiple alignment computed by **MSA** using the default parameter settings. It has four rows, one per sequence. In this case each input sequence has 30 characters, and the output alignment has 33 columns. In general, the input sequences can be of varying lengths. The four sequences in this figure are the first 30 characters of each of the four globin sequences HAHU, HBHU, MYHU, and IGLOB from the globin data set in McClure *et al.* (1994). The default setting in **MSA** does not penalize gaps at the end of the sequence, which explains why the computed alignment puts most of the gaps at the ends.

In the simplest cost model there is a cost function $\text{sub}: \Sigma' \times \Sigma' \rightarrow \mathbb{N}$, such that $\text{sub}(a, b)$ is the cost of substituting a b in the second sequence for an a in the first sequence; also, $\text{sub}(-, b)$ is the cost for columns where the first sequence has a gap and the second has a b , and $\text{sub}(a, -)$ is the cost for columns where the first sequence has an a and the second has a $-$. Function sub should be symmetric, and $\text{sub}(-, -)$ should have value 0. The cost of a pairwise alignment $A_{i,j}$ induced in a rectangular array A of width w is

$$c(A_{i,j}) := \sum_{1 \leq k \leq w} \text{sub}(A[i][k], A[j][k])$$

A crucial point is that the cost of the pairwise alignment $A_{i,j}$ induced by A will always be at least as large as the cost of an optimal alignment of S_i, S_j alone. **MSA** attempts to minimize pairwise alignment costs. This is consistent with the use of Dijkstra’s algorithm in **MSA**, but some other papers on multiple sequence alignment would say that higher values of sub are preferred.

By default, **MSA** supports a more sophisticated cost model where gaps that span several consecutive columns cost less than the sum of sub applied individually to each column. In particular, computing the substitution cost of a column depends on whether the previous column has a $-$ or not. With an optional flag, **MSA** supports a variation on gap costs where gaps at either end of the string cost nothing; gaps at either end are sometimes called *terminal gaps*. We use the notation $c(A_{i,j})$ for all three of the cost models, and rely on the context to specify the cost model.

The basic *sum of pairs* multiple sequence alignment problem is to minimize the pairwise sum

$$c(A) := \sum_{i < j} c(A_{i,j})$$

MSA can solve the sum of pairs multiple sequence alignment, if the user supplies the appropriate flags. By default, however, the **MSA** program solves a variation where each pairwise alignment cost is multiplied by a weight, denoted by $\text{scale}(S_i, S_j)$, with weights chosen from an evolutionary tree of the input sequences. A method described in Altschul *et al.* (1989) is used to compute the pairwise weights. The weights are useful to diminish representational bias caused by having more sequences from some species and fewer sequences from others. We have not modified the algorithms that compute the pairwise weights because they are not space- or time-intensive.

To keep resource usage down **MSA** searches only among alignments whose cost is $\leq U$, for some value of U . The main part of **MSA** assumes that a finite value for U is known. If U is too small, the program does not find any feasible solution. For this reason, it may take several runs, with increasing values of U to even get an alignment, optimal or not. The reason not to use an arbitrarily large value of U initially is that a larger U requires more space and time for program execution.

The value of U used in **MSA** is computed as follows. Define a lower bound L as the weighted sum

$$L := \sum_{i < j} d(S_i, S_j) \cdot \text{scale}(S_i, S_j)$$

of costs of optimal pairwise alignments, where $\text{scale}(S_i, S_j)$ is the weight assigned to the pair of sequences S_i, S_j . This value of L is a lower bound on the cost of a multiple sequence alignment because it assumes that each pair of sequences is aligned optimally, independent of the other sequences. In **MSA**, the difference $U - L$ is stored in the variable **DELTA**, which we denote here by δ . From here, there are three ways to get to U . In each case a different δ is obtained and U is computed as $L + \delta$. One option is for the user to specify δ .

Second, if the user does not provide δ , a preprocessor program computes a heuristic multiple sequence alignment from a star-tree of pairwise alignments in which places where all sequences agree are forced into alignment, and an incremental heuristic is used in the intervals between forced columns. This heuristic is not time- or space-intensive. The heuristic multiple alignment induces a pairwise cost; the difference between that cost and the optimal pairwise cost $d(S_i, S_j)$ is denoted by $\epsilon_{i,j}$.

Then δ is computed by

$$\delta := \sum_{i < j} [\min(\text{maxepsilon}, \epsilon_{i,j}) \cdot \text{scale}(S_i, S_j)]$$

The number **maxepsilon** is set by default to 50, but can be changed by changing one constant in the code.

The third option to get δ is that the user may suppress the preprocessor program by supplying a file with values for $\epsilon_{i,j}$, in which case these values are used in the same way as the computed ϵ values in the second option. Again, this method may yield an upper bound U that is too low.

As Carrillo and Lipman (1988) realized, to solve the sum-of-pairs alignment problem, not all alignments A have to be considered. Define $d(S_1, S_2, \dots, S_K) := \min_A c(A)$. Let L be the sum of pairs lower bound defined as above. Let U be an upper bound on $d(S_1, \dots, S_K)$, and let \mathcal{A} be a multiple sequence alignment of minimal cost. Then for any distinct p and q ,

$$U - L \geq \sum_{i,j} \text{scale}(S_i, S_j) [c(\mathcal{A}_{i,j}) - d(S_i, S_j)] \geq \text{scale}(S_p, S_q) [c(\mathcal{A}_{p,q}) - d(S_p, S_q)]$$

since for all i, j , $c(\mathcal{A}_{i,j}) \geq d(S_i, S_j)$. Rearranging, we derive the Carrillo and Lipman bound

$$\text{scale}(S_i, S_j) c(\mathcal{A}_{i,j}) \leq \text{scale}(S_i, S_j) d(S_i, S_j) + U - L \quad (1)$$

Note that $\Omega(K^2)$ terms are thrown away to arrive at (1); while it is possible for the equation to achieve equality, for K equally dissimilar sequences, $c(\mathcal{A}_{i,j}) - d(S_i, S_j)$ will be overestimated by a factor $O(K^2)$. In the implementation, **MSA** applies (1) only during an initial preprocessing phase.

In the classic dynamic programming algorithm for sequence alignment, a directed acyclic graph is constructed in which source-to-sink paths of weight C correspond to alignments of cost C . For multiple sequence alignment, this graph is conveniently represented with vertices embedded at integer coordinates in K -space and sequences labeling the coordinate axes. A path P starting at the source vertex $s = (0, \dots, 0)$ that ends at vertex $t = (n_1, \dots, n_K)$ represents an alignment of the K prefixes $S_i[1 \dots n_i]$ for $1 \leq i \leq K$. There is a one-to-one correspondence between edges in the path and columns in the alignment. For example, in the alignment of Figure 1 the first two edges are

$$(0, 0, 0, 0) \rightarrow (0, 0, 0, 1) \rightarrow (0, 1, 0, 2)$$

For column c , if we do not have a dash in row i , then the c th edge in the path advances the coordinate in dimension i by 1; if we have a dash in row i , the coordinate in dimension i stays the same.

The algorithm used to compute the shortest path from s to t is a variant of Dijkstra's algorithm (1959). In Dijkstra's algorithm, as used in the next section, each vertex v has a label $v.D$, which is the current cost of a shortest path from s to v . Each edge $e = v \rightarrow w$ has a cost, denoted by $\text{cost}(e)$. When the algorithm explores an edge $v \rightarrow w$, it compares the current $w.D$ (representing a known path to w) against $v.D + \text{cost}(e)$ (representing a path to w through v). The new label $w.D$ is the lesser of the two alternatives.

The algorithm always chooses v to have minimum label D among vertices with unused outgoing edges. The vertices with unexplored outgoing edges are kept in a priority queue, so that v can be found quickly.

Inequality (1) restricts consideration to those paths whose projections¹ onto the planes defined by all pairs S_i and S_j have weight at most $d(S_i, S_j) + U - L$. **MSA** 1.0 also prunes vertices by lower-bounding costs from the current vertex to the sink. If the minimum cost to reach v from the source plus a lower bound on the cost to reach the sink from v is higher than U , v can be pruned. We found that not performing the lower-bound pruning eliminates one field in a record, saving enough space that our new version actually runs faster without the pruning.

For each pair of sequences S_i, S_j , **MSA** computes the standard two-dimensional dynamic programming graph $D_{i,j}$ and, for each vertex, the cost of an optimal alignment passing through that vertex. We call vertices whose cost in $D_{i,j}$ is at most $d(S_i, S_j) + \epsilon_{i,j}$ *admissible* with respect to the pair (i, j) . Returning to the dynamic programming graph in K -space, **MSA** considers only vertices that are admissible with respect to every pair. Admissibility depends only on the pairwise alignment and not the actual costs in the K -dimensional dynamic programming graph. By manually increasing the value of $\epsilon_{i,j}$, the user can improve the chances that **MSA** considers enough of the dynamic programming graph, so that the output alignment is optimal or near-optimal. However, we emphasize that it is generally impractical to set each $\epsilon_{i,j}$ to $U - L$, which is what would be needed to guarantee optimality. Therefore, in practice, **MSA** rarely finds a guaranteed optimal alignment.

FIRST IMPLEMENTATION WITHOUT GAP PENALTIES

We design an algorithm that incorporates inequality (1) and prunes vertices by lower-bounding costs, where pairwise alignments are scored as the sum of single-symbol insertion, deletion, and substitution costs; this algorithm is later extended for an affine gap cost function. Input consists of the sequences S_1, S_2, \dots, S_K of lengths N_1, N_2, \dots, N_K , and the difference $\delta := U - L$. We will write N for $\max_i N_i$.

The algorithm is naturally organized in two stages. In the first stage, $O(K^2)$ pairwise sequence comparisons are performed to determine, for each plane defined by sequence pair S_i and S_j , the set of points $\langle p, q \rangle$ on a path from $\langle 0, 0 \rangle$ to $\langle N_i, N_j \rangle$ of weight at most $d(S_i, S_j) + \epsilon_{i,j}$. Let $F_{i,j}$ be the set of such points $\langle p, q \rangle$ on the face formed by S_i and S_j . At the conclusion of this stage, the $F_{i,j}$ are determined, and the lower bound $L = \sum_{i < j} d(S_i, S_j)$ and upper bound $U = L + \delta$ can be computed. The first stage is neither time- nor space-intensive, so we focus our complete attention on the second stage.

In the second stage, an alignment corresponding to a shortest source to sink path is computed; it is guaranteed that the projections of the path lie wholly within the $F_{i,j}$. Formally the weighted directed acyclic graph $G = \langle V, E \rangle$ has vertex set

$$V := [0, N_1] \times [0, N_2] \times \dots \times [0, N_K]$$

and edge set

$$E := \{v \rightarrow w \mid w - v \in \{0, 1\}^K - \{0\}^K\}$$

where we use the standard componentwise subtraction of vectors. For character a , define $a^1 := a$ and $a^0 := '-'$. Using this notation, edge $\langle v_1, \dots, v_K \rangle \xrightarrow{e} \langle w_1, \dots, w_K \rangle$ has cost

$$\text{cost}(e) = \sum_{i < j} \text{sub}[(S_i[w_i])^{w_i - v_i}, (S_j[w_j])^{w_j - v_j}]$$

Let $R_{i,j}$ be the region in K -space of points $\langle p_1, p_2, \dots, p_K \rangle$ such that $\langle p_i, p_j \rangle \in F_{i,j}$, and define $R := \bigcap_{i < j} R_{i,j}$. Then the paths of interest lie within $G \setminus R$, the subgraph of G induced on the vertex subset R .

There are some subtleties to the shortest-path algorithm, as vertices and edges are generated dynamically in a graph with a mesh vertex structure. We say a vertex or edge “exists” when the algorithm has allocated space for it and has not freed that space. Consider the computation when it examines the edges leaving

¹The projection of the path $\langle p_{11}, p_{12}, \dots, p_{1K} \rangle, \langle p_{21}, p_{22}, \dots, p_{2K} \rangle, \dots, \langle p_{n1}, p_{n2}, \dots, p_{nK} \rangle$, onto the plane i, j is the two-dimensional path $\langle p_{1i}, p_{1j} \rangle, \langle p_{2i}, p_{2j} \rangle, \dots, \langle p_{ni}, p_{nj} \rangle$.

a vertex v with coordinates given by point p . Adjacent vertices are those with coordinates q such that $q - p \in \{0, 1\}^K - \{0\}^K$. Having generated the coordinates of point q for an adjacent vertex, we must find the vertex w corresponding to point q if it exists, or create w if it does not exist. (Vertex w may already exist if another path to point q has already been explored.) We can accomplish this with a trie defined on point coordinates of existing vertices. In the trie, coordinates of a point q spell a path from the root to a leaf, with the leaf pointing to the corresponding vertex w . Leaves all have depth K . Each level corresponds to a dimension.

This algorithm is formalized in Figure 2. Two fundamental data structures are used. Function **PRIORITY()** creates an empty priority queue; **INSERT**(x, k, q) inserts item x with key k into priority queue q ; **EXTRACT**(q) returns and removes the item of minimum key in priority queue q ; **DECREASE**(x, k, q) decreases to k the key of x in q . Function **TRIE()** creates an empty trie; **INSTALL**(x, s, t) installs item x with string s into trie t ; **LOOKUP**(s, t) returns the item with string s in trie t , or Λ if none exists. In addition, two new data types are used. Type **point** is a K -tuple of integers; a point variable p has named fields $p.1, p.2, \dots, p.K$, and is created by $p \leftarrow \text{POINT}(p_1, p_2, \dots, p_K)$. Type **vertex** has two named fields: a distance D , and a point P ; statement $v \leftarrow \text{VERTEX}(d, p)$ creates a vertex v with $v.D = d$ and $v.P = p$. The priority queue of vertices is ranked by distance, while the trie of vertices is constructed over the integer sequences forming

```

algorithm MSA(sequence  $S_1, S_2, \dots, S_K$  ; integer  $\delta$ )
  trie  $T$  ; priority queue  $Q$  ; vertex  $s, t, v, w$  ; point  $p, q$ 
  (Stage 1.)
   $L \leftarrow 0$ 
  for  $I \leftarrow 1$  to  $K - 1$  do
    for  $J \leftarrow I + 1$  to  $K$  do
      Let  $a_1 a_2 \dots a_n = S_I$  and  $b_1 b_2 \dots b_m = S_J$ .
      for  $i \leftarrow 1$  to  $n$  do
        for  $j \leftarrow 1$  to  $m$  do
          Compute  $C_{f,I,J}[i, j] := d(a_1 \dots a_i, b_1 \dots b_j)$  and  $C_{r,I,J}[i, j] := d(a_{i+1} \dots a_n, b_{j+1} \dots b_m)$ .
        od
      od
      for  $i \leftarrow 1$  to  $n$  do
         $F_{I,J}[i] \leftarrow \{j \mid C_{f,I,J}[i, j] + C_{r,I,J}[i, j] \leq C_{f,I,J}[n, m] + \epsilon_{i,j}\}$ 
      od
       $L \leftarrow L + C_{f,I,J}[n, m]$ 
    od
  od
   $U \leftarrow L + \delta$ 
  (Stage 2.)
   $s \leftarrow \text{VERTEX}(0, \text{POINT}(0, 0, \dots, 0))$  ;  $t \leftarrow \text{VERTEX}(\infty, \text{POINT}(N_1, N_2, \dots, N_K))$ 
   $T \leftarrow \text{TRIE}()$  ; INSTALL( $s, s.P, T$ ) ; INSTALL( $t, t.P, T$ )
   $Q \leftarrow \text{PRIORITY}()$  ; INSERT( $s, s.D, Q$ )
  while  $\neg \text{EMPTY}(Q)$  do
     $v \leftarrow \text{EXTRACT}(Q)$  ;  $p \leftarrow v.P$ 
    if  $v = t$  then
      output TRACEBACK( $t$ ) ; return
    fi
    if  $v.D + \sum_{i < j} (\text{scale}(S_i, S_j) \cdot C_{r,i,j}[p.i, p.j]) \leq U$  then
      for all  $q \leftarrow \text{POINT}(q_1, q_2, \dots, q_K)$  s.t.  $(q - p \in \{0, 1\}^K - \{0\}^K)$  and for all  $i < j, q_j \in F_{i,j}[q_i]$  do
         $w \leftarrow \text{LOOKUP}(q, T)$ 
        if  $w = \Lambda$  then
           $w \leftarrow \text{VERTEX}(\infty, q)$  ; INSERT( $w, w.D, Q$ ) ; INSTALL( $w, w.P, T$ )
        fi
        if  $v.D + \text{cost}(p \xrightarrow{e} q) < w.D$  then
           $w.D \leftarrow v.D + \text{cost}(e)$  ; DECREASE( $w, w.D, Q$ )
        fi
      od
    fi
  od
  output There does not exist a multiple sequence alignment with cost at most  $U$ .
end

```

FIG. 2. Algorithm without gap costs.

points. Shortest paths are computed using the distance field; adjacent vertices are generated and accessed using the point field.

Function `TRACEBACK(v)` returns an alignment by tracing backward through the graph from vertex v . (By examining the distance and point fields of preceding vertices, a shortest path can be reconstructed.)

One optimization incorporated in the algorithm of Figure 2 is that edges leaving some vertices may be ignored. When a vertex v is removed from the priority queue, $v.D$ is the length of a shortest path from s to v . It is assumed that all insertion, deletion, and substitution costs—and hence all edge weights—are nonnegative because of the implementation of the priority queue. A lower bound on the length of a path from v to t is $\sum_{i < j} (\text{scale}(S_i, S_j) C_{r,i,j}[v.P.i, v.P.j])$ whose terms are computed in Stage 1. If this quantity plus $v.D$ exceeds U , then the length of any path from s to t through v exceeds U , and edges leaving v do not need to be examined.

The algorithm of Figure 2 uses two types of pruning of vertices. One type, which we call return-cost pruning, eliminates a vertex when the length of a shortest path from the source to the vertex, plus the lower bound on the length of a shortest path from the vertex to the sink, is at least the upper bound U . The second type, which we call Carrillo–Lipman pruning, eliminates a vertex when its projection falls outside the Carrillo–Lipman region on some face. It is natural to ask whether return-cost pruning is stronger than Carrillo–Lipman pruning. The answer is yes, though we do not prove it here; the argument is nearly identical to the derivation of the Carrillo–Lipman bound. One can show that the set of vertices visited by an algorithm that only uses return-cost pruning is always a subset of the set of vertices visited by an algorithm that only uses Carrillo–Lipman pruning. Thus, if one wished to, one could eliminate any reference to the Carrillo–Lipman bound altogether from the algorithm of Figure 2; doing so will never enlarge the part of the dynamic programming graph that is explored to find an optimal alignment. As noted in Results, we found that MSA runs faster without the return-cost pruning due to reducing the space usage, so it makes some sense to leave the Carrillo–Lipman pruning in.

We now give some more details about the concrete implementation of the data structures. Making the realistic assumption that edge costs are integers, we can implement a discrete priority queue with buckets for the possible path lengths, the number of buckets being bounded by U . Then `INSERT()` and `DECREASE()` are constant-time operations, and all `EXTRACT()` operations take $O(U + |R|)$ total time, which is $O(K^2 N + |R|)$. As a minor concern, the statement `VERTEX(∞, \dots)` can be implemented as `VERTEX($U + 1, \dots$)`, since U is a bound on vertex distances.

Since the trie is constructed over integer sequences, the children of a node are conveniently indexed through an array of pointers. Consider accessing the vertex in the trie associated with point $\langle p_1, p_2, \dots, p_K \rangle$ in R . After having traversed the trie on prefix $p_1 p_2 \dots p_{j-1}$, the next coordinate p_j must be in the set $\cap_{i < j} F_{i,j}[p_i]$, since the point is in R . Let m and n be the minimum and maximum integers in this set. Then pointers to child nodes, and an encoding of the intersection, can be represented in a vector of length $O(n - m)$. These vectors tend to be short, and once allocated, never need change in length as new vertices are installed. Generation of adjacent vertices is often short-circuited by “falling off” the trie when there are no vertices with a certain prefix of coordinate values. To minimize space consumption, the P vertex field is not represented as a K -tuple of integers; $v.P$ points to the trie leaf associated with vertex v , and coordinates are recovered from trie node labels by following father links to the root.

One other implementation issue concerns the treatment of holes in the region R . Recall $F_{i,j}[i] := \{j \mid C_{f,i,j}[i, j] + C_{r,i,j}[i, j] \leq d(S_i, S_j) + \epsilon_{i,j}\}$. These sets are implemented as a vector of j values, which can require $O(N)$ space for each set. Alternately, one could approximate the set in $O(1)$ space by storing the minimum and maximum integer elements, on the assumption that the sets are dense, and ignore holes.

ADDING GAP PENALTIES

In this section we explain how to modify the shortest-paths algorithm to accommodate affine gap costs in the alignment scoring function. The pseudocode in this section is a fairly accurate description of the actual code in MSA version 1.0.

When costs are determined character-by-character without special treatment of multicharacter gaps, the shortest paths in the graph representing least-cost alignments satisfy an important property, often called the “principle of optimality.” Specifically, any least-cost path from the source s to vertex w , ending in

```

algorithm COST(edge f, edge e)
  point p, q, r ; integer  $\Delta_e[1 \dots K], \Delta_f[1 \dots K]$  ; character subchar[1 ... K]
  p ← e.tail.P
  q ← f.tail.P
  r ← f.head.P
  for I ← 1 to K do
     $\Delta_e[I] \leftarrow q.I - p.I$ 
     $\Delta_f[I] \leftarrow r.I - q.I$ 
    if  $\Delta_f[I] = 1$  then
      subchar[I] ← S[I][r.I]
    else
      subchar[I] ← DASH()
    fi
  od
  totalcost ← 0
  for I ← 2 to K do
    for J ← 1 to I - 1 do
      if DoNotPenalizeTerminalGaps and (q.I = 0 or q.I = N[I] or q.J = 0 or q.J = N[J]) then
        totalcost ← totalcost + scale(SJ, SI) * sub(subchar[J], subchar[I])
      else
        totalcost ← totalcost + scale(SJ, SI) * (sub(subchar[J], subchar[I]) +
          GAPPENALTY( $\Delta_e[I], \Delta_e[J], \Delta_f[I], \Delta_f[J]$ ))
      fi
    od
  od
  return totalcost
end

```

FIG. 3. Edge cost algorithm.

edge $v \rightarrow w$, must start with a least-cost subpath from s to v . This justifies computing the distance from s to w , denoted $w.D$, by

$$w.D \leftarrow \min_v \{v.D + \text{cost}(v \rightarrow w)\}$$

which the algorithm of the previous section does implicitly.

With affine gap costs, in which the cost of a multicharacter gap in an induced pairwise alignment is a per-character penalty times the length of the gap, plus a gap startup penalty, for initiating a multicharacter gap, the principle of optimality does not apply in the above form. It may be preferable to reach w through a more costly path from s to v that ends with a gap already initiated, so that following with edge $v \rightarrow w$ does not incur additional gap startup cost.

In general, gap startup penalties under the sum-of-pairs scoring function cannot be determined column-by-column. Determining whether a column starts a gap in a pairwise alignment may require knowing the gapping structure of an arbitrary number of previous columns. MSA compromises, and uses a practical scheme of Altschul (1989) for counting gap startup events. In this scheme, which Altschul calls “quasi-natural” gap costs, gap startup events for a given column are determined solely from the preceding column in the alignment. The number of events counted by this scheme is never less than the true number of events under the sum-of-pairs measure; hence the lower bounds used by the shortest-path algorithm for pruning remain valid.

Thus, with affine gap penalties, the cost of an edge f in a path is a function of f and the preceding edge e on the path. The way the costs are computed in MSA 1.0 is shown in Figure 3. Suppose that $e = u \rightarrow v$ and $f = v \rightarrow w$. Let the points p, q, r in K -dimensional space represent the vertices u, v, w , respectively. We compute the cost of edge f under the assumption that it is preceded by edge e . From points p and q , we can determine which sequences already have a dash from edge e : the S_I for which $q.I = p.I$. Otherwise, $q.I - p.I = 1$, in which case S_I contains a real character in the column corresponding to edge e . The difference $q.I - p.I$ is stored in Δ_e ; similarly we compute Δ_f for the edge f . The double loop on I, J at the bottom of the cost algorithm computes the sum of pairwise contributions from edge f . When there are no gaps in the sequences, the pairwise contribution depends only on sub($S_I[r.I], S_J[r.J]$). When there are gaps in the sequences, the GAPPENALTY term, which implements the scheme of Altschul (1989), assigns different costs depending on whether following edge e by edge f is counted as beginning a new gap or continuing an old gap.


```

algorithm MSA(sequence  $S_1, S_2, \dots, S_K$  ; integer  $\delta$ )
  trie  $T$  ; priority queue  $Q$  ; vertex  $s, t, v, w$  ; point  $p, q, r$  ; edge list  $E$  ; edge  $e, f$ 
  (Stage 2.)
   $s \leftarrow \text{VERTEX}(\text{POINT}(0, 0, \dots, 0))$  ;  $t \leftarrow \text{VERTEX}(\text{POINT}(N_1, N_2, \dots, N_K))$ 
   $T \leftarrow \text{TRIE}()$  ;  $\text{INSTALL}(s, s.P, T)$  ;  $\text{INSTALL}(t, t.P, T)$ 
   $e \leftarrow \text{EDGE}(\text{NULL}, s, 0)$ 
   $E \leftarrow \text{EDGELIST}()$  ;  $\text{INCLUDE}(e, E)$ 
   $Q \leftarrow \text{PRIORITY}()$  ;  $\text{INSERT}(e, e.D, Q)$ 
  while  $\neg \text{EMPTY}(Q)$  do
     $e \leftarrow \text{EXTRACT}(Q)$  ;  $v \leftarrow e.\text{head}$ 
     $q \leftarrow v.P$  ;  $p \leftarrow e.\text{tail}.P$ 
    if  $v = t$  then
      output TRACEBACK( $t$ ) ; return
    fi
    if  $e.D + \sum_{i < j} (\text{scale}(S_i, S_j) \cdot C_{r,i,j}[q.i, q.j]) \leq U$  then
      for all  $r \leftarrow \text{POINT}(r_1, r_2, \dots, r_K)$  s.t.  $(r - q \in \{0, 1\}^K - \{0\}^K)$  and for all  $i < j$   $r_j \in F_{i,j}[r_i]$  do
         $w \leftarrow \text{LOOKUP}(r, T)$ 
        if  $w = \text{NULL}$  then
           $w \leftarrow \text{VERTEX}(r)$  ;  $\text{INSTALL}(w, w.P, T)$ 
        fi
         $f \leftarrow \text{FIND}(q, r, E)$ 
        if  $f = \Gamma$  then
           $f \leftarrow \text{EDGE}(v, w, \infty)$  ;  $\text{INCLUDE}(f, E)$  ;  $\text{INSERT}(f, f.D, Q)$ 
        fi
        if  $e.D + \text{COST}(f, e) < f.D$  then
           $f.D \leftarrow e.D + \text{COST}(f, e)$  ;  $\text{DECREASE}(f, f.D, Q)$ 
        fi
      od
    fi
  od
  output There does not exist a multiple sequence alignment with cost at most  $U$ .
end

```

FIG. 4. Modified Stage 2 for affine gap costs.

Figure 4 shows the modifications to the alignment algorithm for affine gap costs. Stage 1 does not change significantly, so we show only the new Stage 2. It is natural to modify Stage 2 so that it is an edge-based version of Dijkstra's algorithm. Instead of computing the shortest path starting at s and ending at each vertex, we compute the shortest path starting at s and ending at each edge. To make things work completely in terms of edges, a dummy source node NULL is introduced and an edge from NULL to s is created to start the algorithm. Now the D distance fields are associated with edges and $e.D$ is the current minimum cost of a path starting at s and ending with edge e .

Looking at the data structures, the priority queue Q now stores edges. The trie still stores vertices to allow lookups by coordinate values. $\text{EDGELIST}()$ is the list of edges. The call $\text{EDGE}(v, w, d)$ makes an edge from vertex v to vertex w with distance label d . $\text{INCLUDE}(e, E)$ adds edge e to edge list E . $\text{FIND}(q, r, E)$ returns the edge with tail q and head r in $\text{EDGELIST } E$ if there is one, or returns Γ if such an edge is not on the list.

Whenever we extract an edge $e = u \rightarrow v$ from the priority queue, we find the points p, q in the mesh graph that the vertices represent. We then look for all points r such that the mesh graph has an edge from the point q to the point r . We either find the vertex w representing r or make a new one and call it w . Let f be the edge $v \rightarrow w$. We compute the cost of f preceded by e and see if this gives a new and/or better way to have a path ending with edge f .

IMPROVEMENTS

In this section we describe the algorithmic changes that we implemented to improve the performance of **MSA**. In planning the improvements, we used two important profiling tools that measure the runtime behavior of C programs. The first tool is the function-level profiler **gprof**, which shows, for each function, how often it is called, from where it is called, and an estimate of the total time it takes over all calls.

The second tool is the basic block-level profiler **tcov**, which shows for each basic block how often it is executed. Both tools can be used to predict how much time will be saved by many types of code changes. **gprof** shows how often memory is allocated, so it is good for predicting how much space will be saved by certain changes. **tcov** shows how often each conditional test succeeds or fails, so it is especially useful in rewriting small intensely used sections of a program.

By profiling **MSA** with **gprof** we immediately learned two important facts:

1. The running time is dominated by three functions:
 - **COST**, which computes the cost of an edge e followed by an edge f
 - **ADJACENT**, which generates the next vertex adjacent to the tail of the edge just extracted from the heap, and
 - **EDGE**, which searches through a small list of edges for an edge between two endpoints, and if the edge is not found creates a record for a new edge and initializes it. Our measurements showed that the vast majority of the time over all calls to **EDGE** was spent searching for edges that already existed.
2. The bulk of the space is taken by records that store edges. This is not surprising since there may be as many as $2^K - 1$ edges emanating from a vertex.

Starting here we naturally set out to investigate exactly when edges are needed. Not surprisingly, we also had to study when vertices are needed. Later we recoded the cost computation to save more time.

We discovered that a major problem with edges is that they were stored in lists at the incoming vertex, i.e., the edge $v \rightarrow w$ would be stored in a list at vertex w . This choice had two virtues: it allowed for easy backtracking from source to sink to report the final alignment, and it was the easiest way to check whether an edge already existed or needed to be created. However, it also had two significant problems: it requires many list searches [which do not take $O(1)$ time] to find records for existing edges, and it precludes several opportunities for deleting edges and vertices. To see the first problem consider what happens each time v is extracted from the priority queue. We will want to find all the edges $v \rightarrow w$, but each such edge is in an arbitrary location on a list at w , and we need to search the list to find it. The second problem is more subtle, and we will explain it later at length. We set out to correct the two major problems, while finding alternate ways to achieve the virtues.

If we instead store the list of outgoing edges $v \rightarrow w$ at v , we can scan the list in $O(1)$ time per edge and we never need to regenerate an edge. We changed the implementation, so that edges are stored in a list at the tail vertex. If the list of outgoing edges is **NULL**, then the vertex has no outgoing edges.

To provide for easy backtracking, we added a **BACKTRACK** field to each edge, which is its preceding edge in some optimal path. This makes backtracking very simple and fast, although it costs one extra pointer per edge. Each edge also has a reference count, which keeps track of how many **BACKTRACK** fields point to it. As we shall see, this reference count is very helpful in detecting when edges and vertices are no longer needed.

With these changes, we can now sketch the proofs of four facts about the creation of vertices and edges. A complete proof would require a line-by-line analysis of the actual code, which would horribly obfuscate the key ideas.

Fact 5.1. Once an edge e is extracted from the priority queue, $e.D$ will never again decrease.

Proof. All edge costs are nonnegative. After the first edge is extracted from the priority queue, whenever an edge h is reinserted into the queue, it must have a preceding edge g such that $g.D \leq h.D$. When $e = w \rightarrow x$ is extracted from the queue, its D field is no bigger than any other D field. The only way $e.D$ could decrease is if we later extract an edge $f = v \rightarrow w$ such that f followed by e provides a better path to get to x . Furthermore, we would have to have $f.D < e.D$. However, f cannot be in the queue when e is extracted because e has a minimum D field, and an edge with a lower D field cannot be inserted into the queue later. \square

Fact 5.2. It is correct to delete (and never recreate) an edge $e = v \rightarrow w$ with $w \neq t$, which has already been extracted from the priority queue, for either of the following two reasons:

1. w has no outgoing edges [can be detected in $O(1)$ time by empty outgoing list], or
2. all edges of the form $w \rightarrow x$ have an edge other than e as backtrack edge [detected in $O(1)$ time by reference count on e].

Proof. Once e is extracted from the queue, all paths that include e followed by an outgoing edge from w are considered. By the previous fact, the cost of such a path can never again decrease. If all outgoing edges of the form $w \rightarrow x$ are eliminated (reason 1), the minimum cost path to x cannot include e .

Any edges that point to e as backtrack edge must be of the form $w \rightarrow x$. The BACKTRACK pointer can be made to point to e only in the phase where e is just extracted from the heap and its succeeding edges are considered. After this time, no new BACKTRACK pointers can be assigned to e . Hence, once we reach a stage where no BACKTRACK pointers point to e (reason 2), the edge e cannot occur in an optimal path and can be deleted. \square

Fact 5.3. If the list of outgoing edges for v ever becomes empty, after v has been visited, then all edges into v can be deleted.

Proof. A key point is that when we first extract an edge whose head is v all outgoing edges from v are created. If all edges coming out of v have been deleted, there is no way to continue a path that gets to v , so no edge coming into v can be in an optimal path. Thus the incoming edges can be deleted. \square

The test in Fact 5.3 is extremely useful in deleting edges, and helps explain why it is a good idea to store the edges at their tail.

Fact 5.4. If vertex v has no outgoing edges or incoming edges after some edge outgoing from v is extracted, then v can never occur on a optimal path. In this case, v can be deleted, and need not be recreated.

Proof. If v is deleted, this is because it either has no incoming edges or no outgoing edges remaining, after having some. This must be because all the edges have themselves been deleted. By Facts 5.2 and 5.3, these edges cannot occur on an optimal path and do not need to be recreated. Hence v cannot occur on an optimal path and need not be recreated. \square

The tests in the preceding facts allow us to aggressively free space used by edges and vertices. It is important to combine the tests for the two kinds of objects, even though most of the space is taken up by edges. This is because the deletion of a vertex permits the deletion of all its remaining outgoing or incoming edges, causing a cascade of deletions.

Next we turn to some improvements that reduce the time spent computing COST. The first improvement we made was to inline COST by moving its code inside the MSA function that has the basic loop for Dijkstra's algorithm. An obvious, but minor advantage of inlining is that the time needed for the actual function call is saved.

A much more significant advantage of inlining COST is that we can use invariants that apply across consecutive calls. Recalling the main loop from the previous section, suppose e has just been extracted from the queue. Let the endpoints of e have coordinate values p and q . We then find each point r that can follow q and call COST with points p, q, r as arguments. The key points are that there will usually be multiple values of r for one p, q pair, but much of the COST computation depends only on p and q . After inlining COST, we identified those pieces of the COST computation that depended only on p, q and pulled them out of the loop on r .

The second collection of improvements involved the fundamental comparison of D labels in Dijkstra's algorithm. Recall that if we just extracted $e = p \rightarrow q$ (using coordinates) from the heap and $f = q \rightarrow r$ is a succeeding edge, we compare whether $e.D + \text{cost}(e, f) < f.D$. Our experiments with `tcov` made it clear that the test fails very often.

Our experiments suggested that we should see if the test would fail before all of the costs in COST are calculated. That is, it may be the case that $e.D + \text{extra} \geq f.D$, where $\text{extra} < \text{cost}(e, f)$. As a start, we tried with $\text{extra} = 0$, avoiding the COST computations entirely. Surprisingly this showed a measurable speedup. Then we noticed that the COST computation, now inlined, had a double loop on i and j , where each loop iteration adds to the cost. A natural place to put a more refined partial test is after each iteration of the outer loop on i , using extra as the pairwise sum of costs for all $j < i$.

Finally, we noticed from a `tcov` profile that there is a different way to arrange the double loop on i and j , which might allow for more early exits with a sufficiently large value of extra . In the original formulation

the first iteration on i did 1 value of j , the second iteration on i did 2 values of j and so on, until the last iteration did $K - 1$ values of j . We reversed the loop on i , so that the first iteration did $K - 1$ values, the second $K - 2$, and so on.

RESULTS

In this section, we compare the performance of **MSA** 1.0 with our improved version on real data sets. We ran all the tests on a Sun SparcStation 10 with 128 Mbytes of RAM. We used the cc C compiler with the -O flag for optimization. To measure time we took the sum of the user and system times reported by the **time** command. To measure space usage we ran the **top** command while **MSA** was running, and noted the peak amount of space usage that it reported.

We used data sets extracted from the data sets presented by McClure *et al.* (1994). They compare 4 data sets consisting of 12 globin sequences, 12 kinase sequences, 12 protease sequences, and 12 RH sequences, respectively. We extracted subsets of these 12-element sets that would give rise to **MSA** runs of moderate time and space usage. This enables us to measure the improvement in resource usage without affecting the system too much. We did do two runs in which **MSA** 1.0 far exceeded the 128 Mbytes of RAM (by using virtual memory), causing thrashing; in one of those two cases we had to stop the run to permit other users to get fair access to the computer.

In all runs, we used **MSA** with the default parameter settings, except that we made δ large enough for **MSA** to find some alignment. As the default settings consider restricted $F_{i,j}$ regions, the alignments produced are not provably optimal.

The data sets we used, in the nomenclature of (McClure *et al.*, 1994) are

- Globins A: HUMA (human hemoglobin, α -chain), HBHU (human hemoglobin β -chain), MYHU (human myoglobin), IGLOB (insect hemoglobin from *Chironomus thummi*), GPUGNI (nonlegume hemoglobin from swamp oak), GGZLB (bacteria, *Viteoscilla* sp.), GPYL (legume, yellow lupine).
- Globins B: HAOR (duckbill platypus, α -chain hemoglobin), HADK (duck, α -chain hemoglobin), HBOR (duckbill platypus, β -chain hemoglobin), HBDK (duck, β -chain hemoglobin, MYHU, MYOR (duckbill platypus, myoglobin), IGLOB, GPYL, GPUGNI, GGZLB.
- Proteases A: HIV-I (human immunodeficiency virus, Type I), SRV-I (simian retrovirus, type I), MoMLV (Moloney murine leukemia virus), 17.6 (retrotransposon from *Drosophila melanogaster*), PEPH (human pepsin sequence).
- Proteases B: PEPH, MoMLV, CaMV (cauliflower mosaic virus), Copia (retrotransposon from *Drosophila melanogaster*).
- Kinases A: CD28 (*S. cerevisiae*), MLCK (rat skeletal muscle), PSKH (HeLa cell), CAPK (bovine cardiac muscle), WEE1 (dual specificity kinase from *S. pombe*).
- Kinases B: CD28, MLCK, PSKH, CAPK, WEE1, CSRC (chicken oncogenic protein).
- Kinases C: CD28, CAPK, WEE1, PDGMR (mouse, PDGF receptor).
- RH A: HTLV-II (human T-cell leukemia virus), RSV (Rous sarcoma virus), MoMLV (Moloney murine leukemia virus), 17.6 (retrotransposon from *Drosophila melanogaster*), HBV (human hepatitis B virus, ayw strain).

Table 1 shows the relative time and space usage of **MSA** 1.0 (old) and **MSA** 2.0 (new) on these data sets. In the case of kinases B, we had to stop the old **MSA** run after it had run overnight and used more than 300 Mbytes of memory. In the case of RH A, the old **MSA** run thrashed extensively, and therefore, the actual wallclock time from start to finish was 44 h (much larger than the cpu time usage reported in the table) compared to approximately 1 h for the new version. We note that the running time seems to depend more on how similar the sequences are, rather than on how many sequences there are. For example, proteases B has only 4 sequences, but takes much longer than proteases A, which has 5 sequences.

Removing the pruning of edges by lower-bound to sink improved the time and space further. For example, it improved globins A to 135 s/4.3 Mbytes, proteases B to 7 min/5.2 Mbytes, and kinases A to 8 min/6.0 Mbytes.

The last column of Table 1 measures how many motifs **MSA** correctly aligned out of those given by McClure *et al.* (1994) for the sequences. Each motif is a block of consecutive alignment columns, ranging

TABLE 1. COMPARISON OF RESOURCE USAGE FOR TWO VERSIONS OF **MSA**^a

# Data set	Old time	New time	Old space Mbytes	New space Mbytes	Correct motifs
Globins A (7)	429 s	157 s	15.1	4.8	4.86/5
Globins B (10)	385 s	130 s	11.3	6.1	5.00/5
Proteases A (5)	79 s	37 s	5.7	3.0	2.80/3
Proteases B (4)	24 min	9 min	81	5.8	0.50/3
Kinases A (5)	35 min	10 min	67	6.6	8.00/8
Kinases B (6)	Stopped	118 m	> 300	25	8.00/8
Kinases C (4)	530 s	210 s	35	4.7	6.75/8
RH A (5)	394 min	68 min	380	31	2.60/4

^aNumber of sequences in each data set is in parentheses.

from one to five columns, that a good alignment should contain, based on biochemical and crystallographic considerations. For each data set, we report the sum over the motifs of the fraction of sequences correctly aligned, followed by the number of motifs in the data set. For example on globins B, with 10 sequences, all five motifs were correctly aligned. On globins A, with 7 sequences, **MSA** aligned four of the five motifs correctly, and in the fifth motif, 6 of the 7 sequences were correctly aligned. Generally, on data sets with more sequences, more motifs were correctly aligned. Since the motifs are defined based on sets of 12 sequences, they may not be as easy to detect in small subsets of those 12.

DISCUSSION

We have made improvements in the space and time usage of the **MSA** program for multiple sequence alignment. Since memory is finite on any computer, our improvements make feasible many problems that were not feasible with the first version of **MSA**. The changes to the code involved what would generally be considered low-level implementation decisions. Our work serves as yet another example of the fact that in programs that allocate a lot of memory, it pays to be careful in looking for opportunities to free memory. Also, in programs that have a tight inner loop that takes most of the time, it pays to try a variety of methods of coding that inner loop.

ACKNOWLEDGMENTS

S.K.G. and A.A.S. thank Dr. Sandhya Dwarkadas for useful discussions about **MSA** and Prof. Willy Zwaenepoel for his encouragement. J.D.K. thanks Prof. Marcella McClure for providing her data sets (McClure *et al.*, 1994). Thanks to the referees for useful suggestions and corrections. Research of S.K.G. was supported by grants from the National Science Foundation and the Texas Advanced Technology Program. Research of J.D.K. was supported by a Department of Energy Human Genome Distinguished Postdoctoral Fellowship. Research of the A.A.S. was supported by a contract from IBM and a grant from the National Institutes of Health. Some of the work of J.D.K. was carried out at the Department of Computer Science of the University of California at Davis.

REFERENCES

- Altschul, S.F. 1989. Gap costs for multiple sequence alignment. *J. Theor. Biol.* 138, 297–309.
- Altschul, S.F., Carrol, R.J., and Lipman, D.J. 1989. Weights for data related by a tree. *J. Mol. Biol.* 207, 647–653.
- Barton, G.J., and Sternberg, M.J.E. 1987a. Evaluation and improvements in the automatic alignment of protein sequences. *J. Mol. Biol.* 198, 327–337.
- Barton, G.J., and Sternberg, M.J.E. 1987b. A strategy for the rapid multiple alignment of protein sequences. *Protein Eng.* 1, 89–94.

- Bodlaender, H., Downey, R.G., Fellows, M.R., and Wareham, H.T. 1994. The parameterized complexity of sequence alignment and consensus. *Proc. 5th Symp. Combinatorial Pattern Matching, Lecture Notes Comp. Sci.* 807, 15–30.
- Carrillo, H., and Lipman, D. 1988. The multiple sequence alignment problem in biology. *SIAM J. Appl. Math* 48, 1073–1082.
- Chan, S.C., Wong, A.K.C., and Chiu, D.K.Y. 1992. A survey of multiple sequence comparison methods. *Bull. Math. Biol.* 54, 563–598.
- Dijkstra, E.W. 1959. A note on two problems in connexion with graphs. *Numerische Mathematik* 1, 269–271.
- Feng, D., and Doolittle, R. 1987. Progressive sequence alignment as a prerequisite to correct phylogenetic trees. *J. Mol. Evol.* 25, 351–360.
- Higgins, D.G., Bleasby, A.J., and Fuchs, R. 1992. Clustal V: Improved software for multiple sequence alignment. *CABIOS* 8, 189–191.
- Kececioğlu, J. 1993. The maximum weight trace problem in multiple sequence alignment. *Proc. 4th Symp. Combinatorial Pattern Matching, Springer-Verlag Lecture Notes Comp. Sci.* 684, 106–119.
- Lipman, D.J., Altschul, S.F., and Kececioğlu, J.D. 1989. A tool for multiple sequence alignment. *Proc. Natl. Acad. Sci. U.S.A.* 86, 4412–4415.
- Maier, D. 1978. The complexity of some problems on subsequences and supersequences. *J. ACM* 25, 322–336.
- McClure, M.A., Vasi, T.K., and Fitch, W.M. 1994. Comparative analysis of multiple protein-sequence alignment methods. *Mol. Biol. Evol.* 11, 571–592.
- Subbiah, S., and Harrison, S.C. 1989. A method for multiple sequence alignment with gaps. *J. Mol. Biol.* 209, 539–548.
- Taylor, W.R. 1987. Multiple sequence alignment by a pairwise algorithm. *CABIOS* 3, 81–87.
- Taylor, W.R. 1988. A flexible method to align large numbers of biological sequences. *J. Mol. Evol.* 28, 161–169.
- Wang, L., and Jiang, T. On the complexity of multiple sequence alignment. *J. Comput. Biol.* 1, 337–348.

Address reprint requests to:
Alejandro A. Schäffer
Department of Computer Science
Rice University
6100 Main Street
Houston, TX 77005-1892
schaffer@cs.rice.edu

Received for publication December 5, 1994; accepted as revised May 22, 1995.